



5 kW campaigns — corrected era (solid) vs voided pre-fix attempt (dashed)

fitness (v12 objective; lower = better,
rejects = 1e9 off-scale)

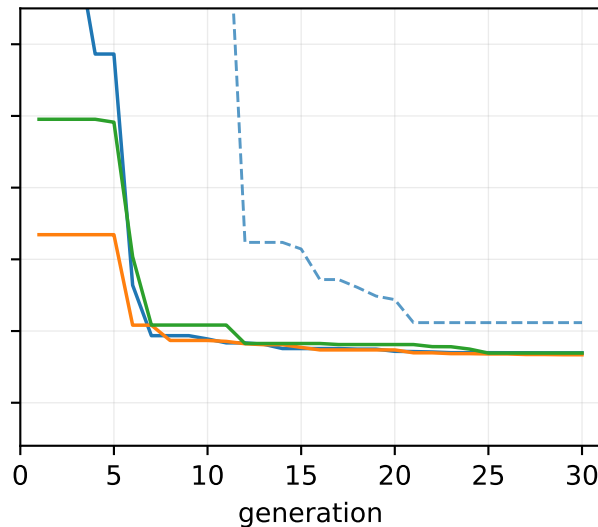
18.0 m



21.2 m



25.0 m



The same three chain lengths, two eras of the evaluator: the corrected campaigns (solid lines) and the earlier attempt, voided after a bug was found (dashed). The traces run almost in step — and at 18 m the voided attempt even scores deeper (-6.66 vs -6.22) — the trap: the pre-fix evaluator missed a sanity check on the hub geometry, its scores were artificially generous, and ‘better’ was an artifact. The voided run is kept as evidence of why the check exists, not as a result.